

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

In [4]:

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

In [5]:

```
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                          n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
```

In [6]:

```
X.shape, y.shape
```

Out[6]:

```
((50000, 15), (50000,))
```

Splitting data into train and test

In [7]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

In [8]:

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[8]:

```
((37500, 15), (37500,), (12500, 15), (12500,))
```

Initialize weights

In [9]:

```
def initialize_weights(row_vector):
    ''' we will initializing our weights and bias'''
    w=np.zeros_like(row_vector)
    b=0
    return w,b
```

In [10]:

```
initialize_weights(X[0])
```

Out[10]:

```
(array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.]), 0)
```

In [11]:

```
dim=X_train[0]
w,b = initialize_weights(X[0])
print('w =', (w))
print('b =', str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Compute sigmoid

In [12]:

```
from math import *
def sigmoid(z):
    ''' we will return sigmoid of z'''
    # compute sigmoid(z) and returnz
    return 1/(1+np.exp(-z))
```

In [13]:

```
def grader_sigmoid(z):
    val=sigmoid(z)
    assert (val==0.8807970779778823)
    return True
grader_sigmoid(2)
```

Out[13]:

True

Compute logloss

In [14]:

```
def logloss(y_true,y_pred):
    loss=-1*np.mean(y_true*np.log10(y_pred)+(1-y_true)*np.log10(1-y_pred))
    return loss
```

In [15]:

```
def grader_logloss(true,pred):
    loss=logloss(true,pred)
    assert (np.round(loss,6)==0.076449)
    return True
true=np.array([1,1,0,1,0])
pred=np.array([0.9,0.8,0.1,0.8,0.2])
grader_logloss(true,pred)
```

Out[15]:

True

Compute gradient w.r.t 'w'

In [16]:

```
def gradient_dw(x,y,w,b,alpha,N):
    ''' we will compute the gardient w.r.to w '''
    dw= ((x*((y-sigmoid(np.dot(w.T,x) + b)))-((alpha/N)*w)))
    return dw
```

In [17]:

```
def grader_dw(x,y,w,b,alpha,N):
    grad_dw=gradient_dw(x,y,w,b,alpha,N)
    assert (np.round(np.sum(grad_dw),5)==4.75684)
    return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
-2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
  3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
-0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
-0.05708987, -0.064768 ,  0.18012332, -0.16880843, -0.27079877])
grad_b=0.5
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[17]:

True

Compute gradient w.r.t 'b'

In [18]:

```
def gradient_db(x,y,w,b):  
    '''In this function, we will compute gradient w.r.to b '''  
    db=y-sigmoid(np.dot(w,x)+b)  
    return db
```

In [19]:

```
def grader_db(x,y,w,b):  
    grad_db=gradient_db(x,y,w,b)  
    assert(np.round(grad_db,4)==-0.3714)  
    return True  
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,  
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,  
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])  
grad_y=0.5  
grad_b=0.1  
grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,  
                -0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,  
                -0.05708987, -0.064768  ,  0.18012332, -0.16880843, -0.27079877])  
alpha=0.0001  
N=len(X_train)  
grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[19]:

True

Logistic Regression

In [21]:

```
def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):  
    ''' we will implementing logistic regression from scratch'''  
    train_loss = []  
    test_loss = []  
    w,b = initialize_weights(X_train[0]) # Initializing the weights  
  
    for i in tqdm(range(epochs)):  
        for k,j in zip(X_train,y_train):  
            dw=gradient_dw(k,j,w,b,alpha,1)  
            db=gradient_db(k,j,w,b)  
            w=w+eta0*dw #updating w, b  
            b=b+eta0*db  
        predicted_train=pred(w,b,X_train)  
        train_loss.append(logloss(y_train,predicted_train)) # storing all the train loss values in a list  
        predicted_test=pred(w,b,X_test)  
        test_loss.append(logloss(y_test,predicted_test)) # storing all the train loss values in a list  
        print("epochs {}".format(i))  
        print("train loss {}".format(train_loss[i]))  
        print("test loss {}".format(test_loss[i]))  
    return w,b,train_loss,test_loss
```

In [22]:

```
from tqdm import tqdm  
alpha=0.0001  
eta0=0.0001  
N=len(X_train)  
epochs=20  
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

5%|██████████| 1/20 [00:01<00:37, 1.98s/it]

epochs 0
train loss 0.1754606247360309
test loss 0.17595770726331647

10%|██████████| 2/20 [00:03<00:35, 2.00s/it]

epochs 1
train loss 0.16867416744464175
test loss 0.16940149097552995

15%|██████████| 3/20 [00:05<00:32, 1.94s/it]

epochs 2 train loss 0.1663936302047816 test loss 0.16720724323524605					
20% ████████████████████			4/20	[00:08<00:33, 2.08s/it]	
epochs 3 train loss 0.1653697408374924 test loss 0.16621787754164866					
25% ██████████████████			5/20	[00:10<00:32, 2.16s/it]	
epochs 4 train loss 0.1648581973924498 test loss 0.1657198487150934					
30% ████████████████████			6/20	[00:12<00:29, 2.08s/it]	
epochs 5 train loss 0.1645890763819657 test loss 0.1654555081976047					
35% ████████████████████			7/20	[00:14<00:26, 2.06s/it]	
epochs 6 train loss 0.16444340426111953 test loss 0.1653110544953986					
40% ████████████████████			8/20	[00:16<00:24, 2.01s/it]	
epochs 7 train loss 0.16436318230671837 test loss 0.16523069542312474					
45% ████████████████████			9/20	[00:18<00:22, 2.01s/it]	
epochs 8 train loss 0.16431849940861387 test loss 0.1651854473080727					
50% ████████████████████			10/20	[00:20<00:20, 2.06s/it]	
epochs 9 train loss 0.16429340952198623 test loss 0.16515973510755802					
55% ████████████████████			11/20	[00:22<00:18, 2.05s/it]	
epochs 10 train loss 0.16427923378687076 test loss 0.16514501157574715					
60% ████████████████████			12/20	[00:24<00:16, 2.02s/it]	
epochs 11 train loss 0.1642711834339083 test loss 0.16513652026883494					
65% ████████████████████			13/20	[00:26<00:13, 2.00s/it]	
epochs 12 train loss 0.16426659068998162 test loss 0.1651315878603367					
70% ████████████████████			14/20	[00:28<00:11, 1.97s/it]	
epochs 13 train loss 0.1642639588605032 test loss 0.16512870040110492					
75% ████████████████████			15/20	[00:30<00:09, 1.97s/it]	
epochs 14 train loss 0.16426244369614623 test loss 0.16512699518042384					
80% ████████████████████			16/20	[00:32<00:08, 1.97s/it]	

```
epochs 15
train loss 0.16426156688862956
test loss 0.16512597788517092
```

```
epochs 16
train loss 0.1642610564389263
test loss 0.16512536379282627
```

```
epochs 17
train loss 0.16426075714134006
test loss 0.16512498801128778
```

```
epochs 18
train loss 0.16426058013656883
test loss 0.1651247544742231
```

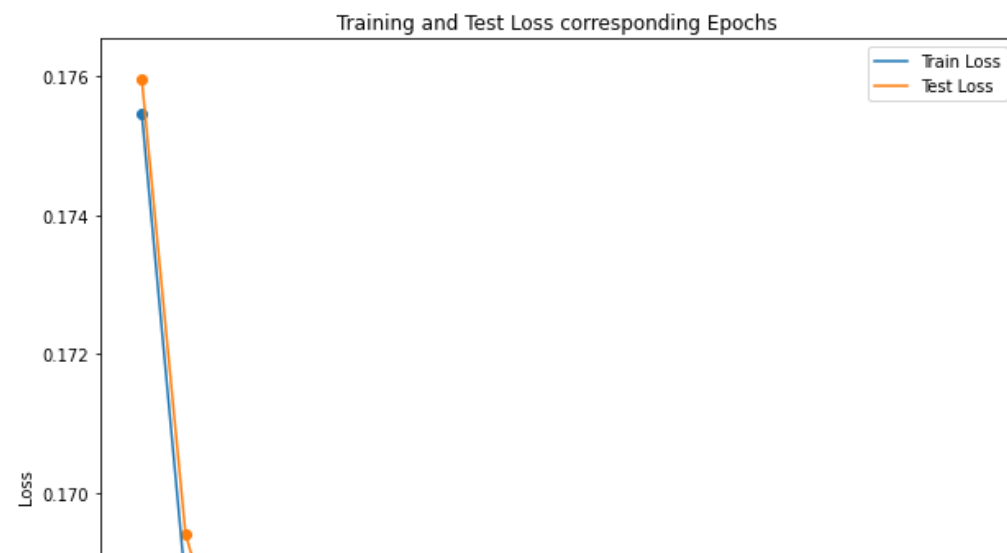
```
epochs 19
train loss 0.1642604743701081
test loss 0.16512460682798147
```


```
print (w)
print (b)
```

```
[-4.29140945e-01  1.92805460e-01 -1.48151251e-01  3.37889601e-01
 -2.20518861e-01  5.69311482e-01 -4.45010836e-01 -8.99621603e-02
  2.21483121e-01  1.73425928e-01  1.98413267e-01 -4.20765000e-04
 -8.10591107e-02  3.38895504e-01  2.29521948e-02]
-0.8895989285186564
```

In [25]:

```
import matplotlib.pyplot as plt
plt.figure(figsize = (10,10))
plt.plot(range(1,epochs+1),train_loss)
plt.plot(range(1,epochs+1),test_loss)
plt.scatter(range(1,epochs+1),train_loss)
plt.scatter(range(1,epochs+1),test_loss)
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training and Test Loss corresponding Epochs")
labels = ["Train Loss" , "Test Loss"]
plt.legend(labels, loc = "upper right")
plt.show()
```





In [32]:

```
def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
            predict.append(1)
        else:
            predict.append(0)
    return np.array(predict)
print("Training Accuracy without Sklearn Implemetation : {0}".format(1-np.sum(y_train - pred(w,b,X_train))/len(X_train)))
print("Test Accuracy Without Sklearn Implementation : {0}".format(1-np.sum(y_test - pred(w,b,X_test))/len(X_test)))
```

Training Accuracy without Sklearn Implemetation : 0.9523733333333333
 Test Accuracy Without Sklearn Implementation : 0.95024